

Lab4: InfoGAN

Department of Computer Science, NCTU

TA: Alan(李懿倫)、Jonson(黃正昕)


Important Date

- Experiment Report Submission Deadline: 12/29 11:59 a.m
- Demo date: 12/29
- Zip all files in one file
 - Report (.pdf)
 - Source code (.py)
- Name it like 「 DLP_LAB4_yourID_name.zip 」
 - ex: DLP_LAB4_409551015_李懿倫.zip
 - Email it to jshuang.cs09@nycu.edu.tw with subject MTK_DLP_LAB4_yourID_name


Lab Objective

- In this lab, you need to implement InfoGAN for generation of handwritten digit images.
- Handwritten digit images generation
 - Noise + fixed latent code -> generate handwritten images of the same digit but different appearance.


e.g. noise + fixed latent code of digit '0' ->



GAN (Generative Adversarial Networks)



I just want to learn generator!



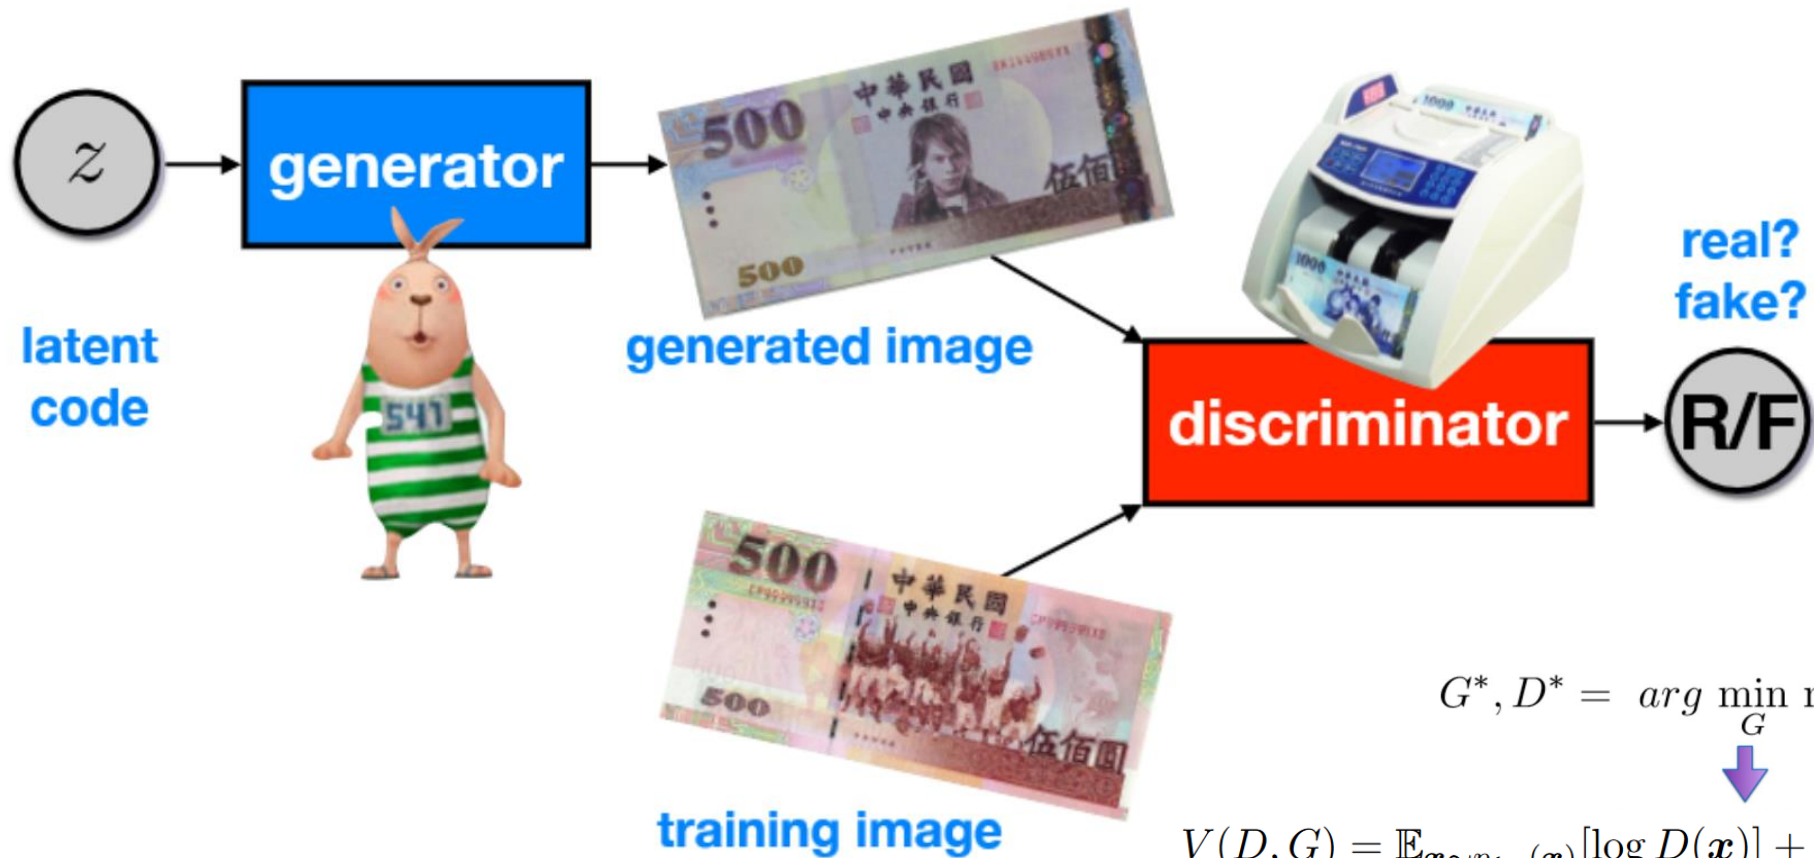
might look like a fake image,
how to get it more realistic?

👉 impose **adversarial loss** on **data distribution**

GAN (Generative Adversarial Networks)

generator: try to generate more realistic images to cheat discriminator

discriminator: try to distinguish whether the image is generated or real



$$G^*, D^* = \arg \min_G \max_D V(D, G)$$



$$V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

Vanilla GAN Loss

Adversarial Loss

$$L_D = -\log(D(x)) - \log(1 - D(G(z)))$$

$$\begin{cases} L_G = \log(1 - D(G(z))) \\ L_G = -\log(D(G(z))) \end{cases}$$

$$G^*, D^* = \arg \min_G \max_D V(D, G)$$



$$V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

GAN-Train the Discriminator

- Minimize $L_D = -\log(D(x)) - \log(1 - D(G(z)))$
 - Passing a batch of real data through D and calculate $\log(D(x))$
 - `nn.BCELoss()(D(x), 1)`

```
## Train D with all-real batch
optimD.zero_grad()
# Get a batch of real data
real_data = data[0].to(device)
b_size = real_data.size(0)
label = torch.full((b_size,), real_label, device=device)
# Forward real batch to D
output = netD(real_data)
# Calculate loss on all-real batch
loss_real = criterionD(output, label)
# Calculate the gradients for real batch
loss_real.backward()
```

GAN-Train the Discriminator

- Minimize $L_D = -\log(D(x)) - \log(1 - D(G(z)))$
 - Sample a noise z and construct a fake data $G(z)$
 - Passing a batch of fake data through D and calculate $\log(1 - D(G(z)))$
 - `nn.BCELoss()(D(G(z)), 0)`

```
## Train D with all-fake batch
# Generate batch of latent vectors
noise = torch.randn(b_size, nz, 1, 1, device=device)
# Generate fake image batch with Generate
fake_data = netG(noise)
label.fill_(fake_label)
# Classify all fake batch with D
output = netD(fake_data.detach())
# Calculate D's loss on the all-fake batch
loss_fake = criterionD(output, label)
# Calculate the gradients for fake batch
loss_fake.backward()
# Update parameters
optimD.step()
```




GAN-Train the Generator

- Minimize $L_G = -\log(D(G(z)))$
 - Passing the fake data through D and calculate $\log(D(G(z)))$
 - `nn.BCELoss()(D(G(z)), 1)`

```
## Train G
optimG.zero_grad()
# label for generator should be real_label -> fooling D
label.fill_(real_label)
# Since we just updated D, perform another forward path of all-fake batch through D
output = netD(fake_data)
# Calculate G's loss based on this output
loss_G = criterionD(output, label)
# Calculate gradient for G
loss_G.backward()
# Update parameters
optimG.step()
```


From GAN to DCGAN

- Replace fully connected layers with convolutions
- Use batch normalization after each layer



GAN

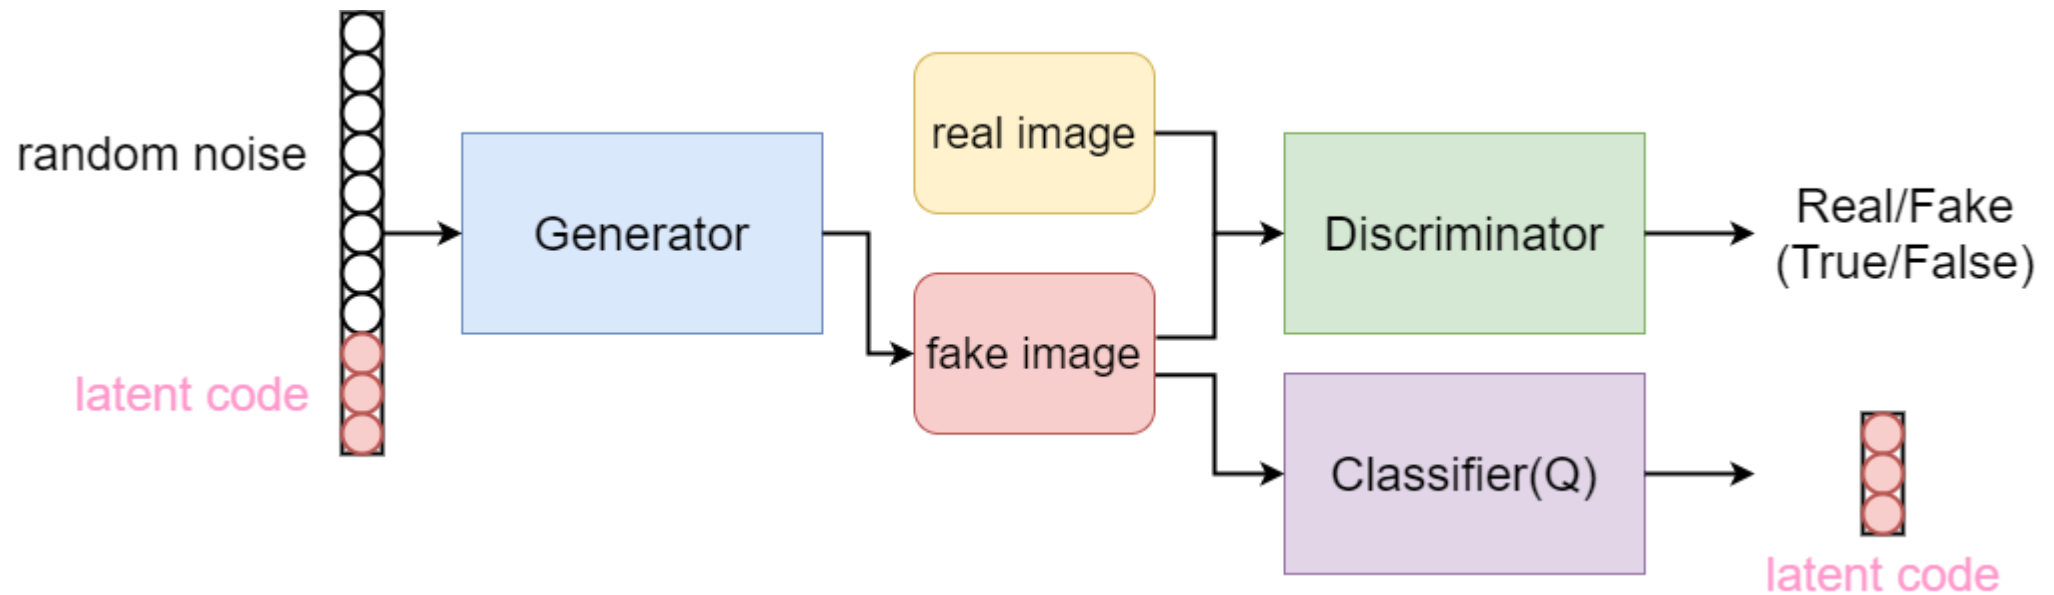
fully connected (project)
from 100 to 16x1024
then reshape to 4x4x1024



DCGAN

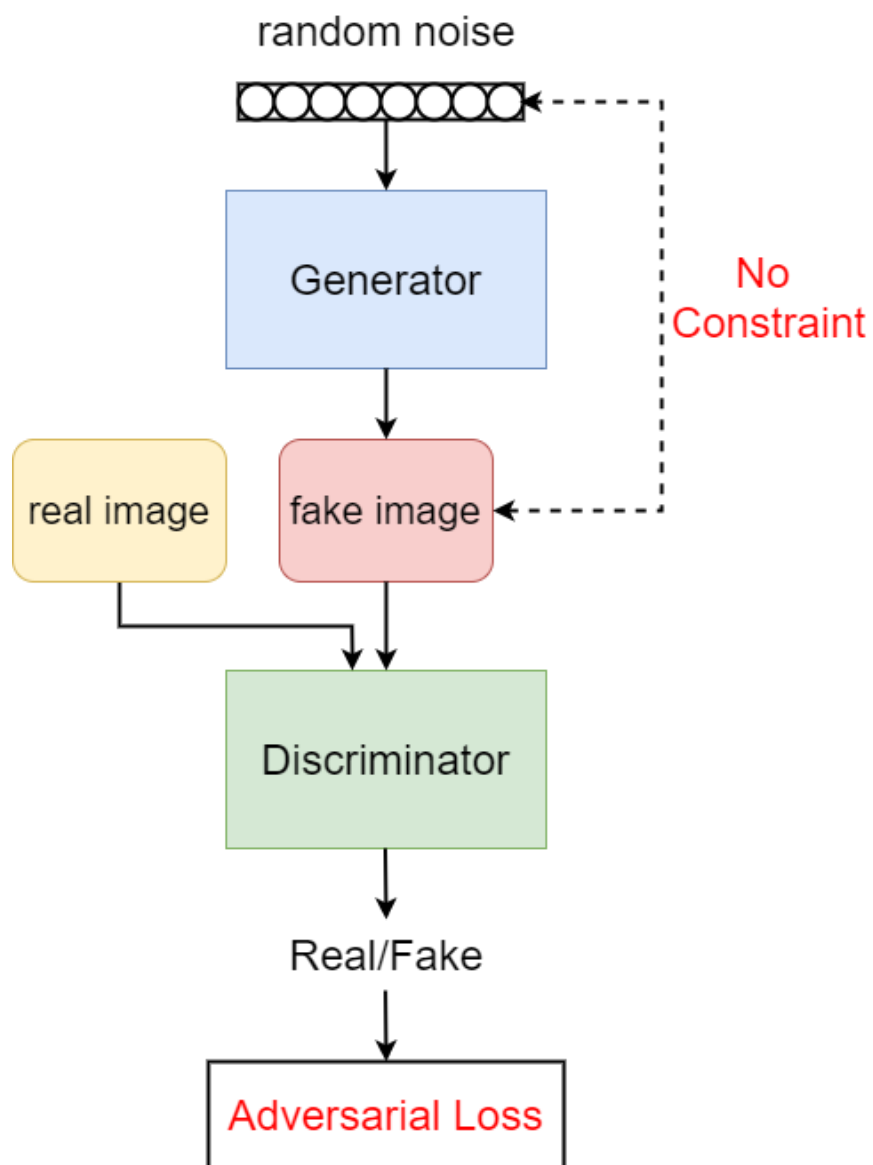
InfoGAN

- InfoGAN: Interpretable Representation Learning by Information Maximizing Generative Adversarial Nets

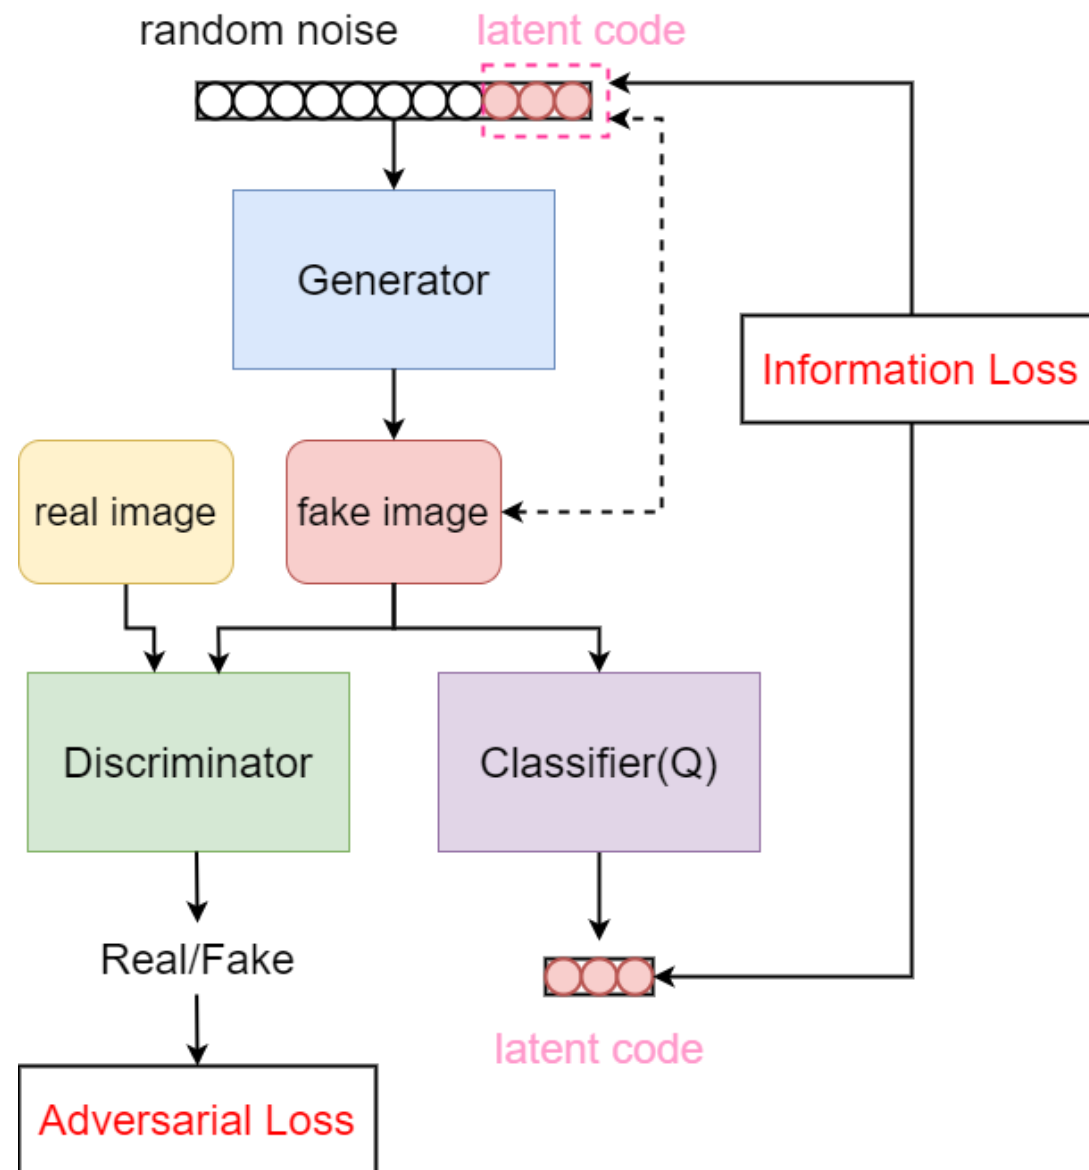


Difference

GAN



InfoGAN



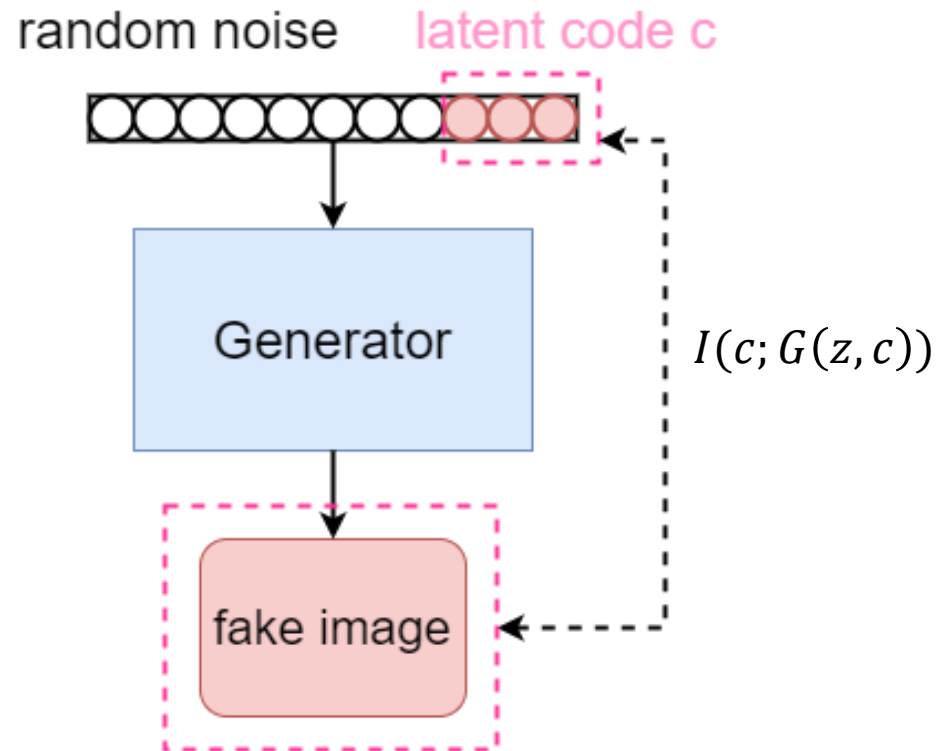

InfoGAN

- Mutual Information:

$$I(X; Y) = H(X) - H(X|Y) = H(Y) - H(Y|X)$$

$$H(Y) = - \sum_y \log p(y)p(y)$$

$$I(X; Y) = H(X) - H(X|Y) = H(Y) - H(Y|X)$$



InfoGAN-prof

$$\begin{aligned} I(c; G(z, c)) &= H(c) - H(c|G(z, c)) \\ &= E_{x \sim G(z, c)} \left[E_{c' \sim P(c|x)} [\log P(c'|x)] \right] + H(c) \\ &= E_{x \sim G(z, c)} \left[D_{KL}(P(\cdot|x) || Q(\cdot|x)) + E_{c' \sim P(c|x)} [\log Q(c'|x)] \right] + H(c) \\ &\geq E_{x \sim G(z, c)} \left[E_{c' \sim P(c|x)} [\log Q(c'|x)] \right] + H(c) \\ &= E_{c \sim P(c), x \sim G(z, c)} [\log Q(c|x)] + H(c) \end{aligned}$$



$$L_I(G, Q) = c \cdot \log Q(G(z, c)) + \boxed{H(c)}$$

fixed value

$$\text{maximize } I(c; G(z, c)) = \text{maximize } L_I(G, Q) = \text{maximize } c \cdot \log Q(G(z, c)) = \boxed{\text{minimize } -c \cdot \log Q(G(z, c))}$$

Cross Entropy Loss

InfoGAN Loss

Adversarial Loss

$$L_D = -\log(D(x)) - \log(1 - D(G(z)))$$


$$\begin{cases} L_G = \log(1 - D(G(z))) \\ L_G = -\log(D(G(z))) \end{cases}$$

Information Loss

$$L_I(Q, G) = -c \cdot \log(Q(G(z, c)))$$

InfoGAN (share-layer)

- In practice, D and Q will share conv layers to share the information.



```
class Discriminator(nn.Module):
    def __init__(self):
        super(Discriminator, self).__init__()

        # shared layer of discriminator
        self.main = nn.Sequential(
            #####
            # To Do
            #####
        )

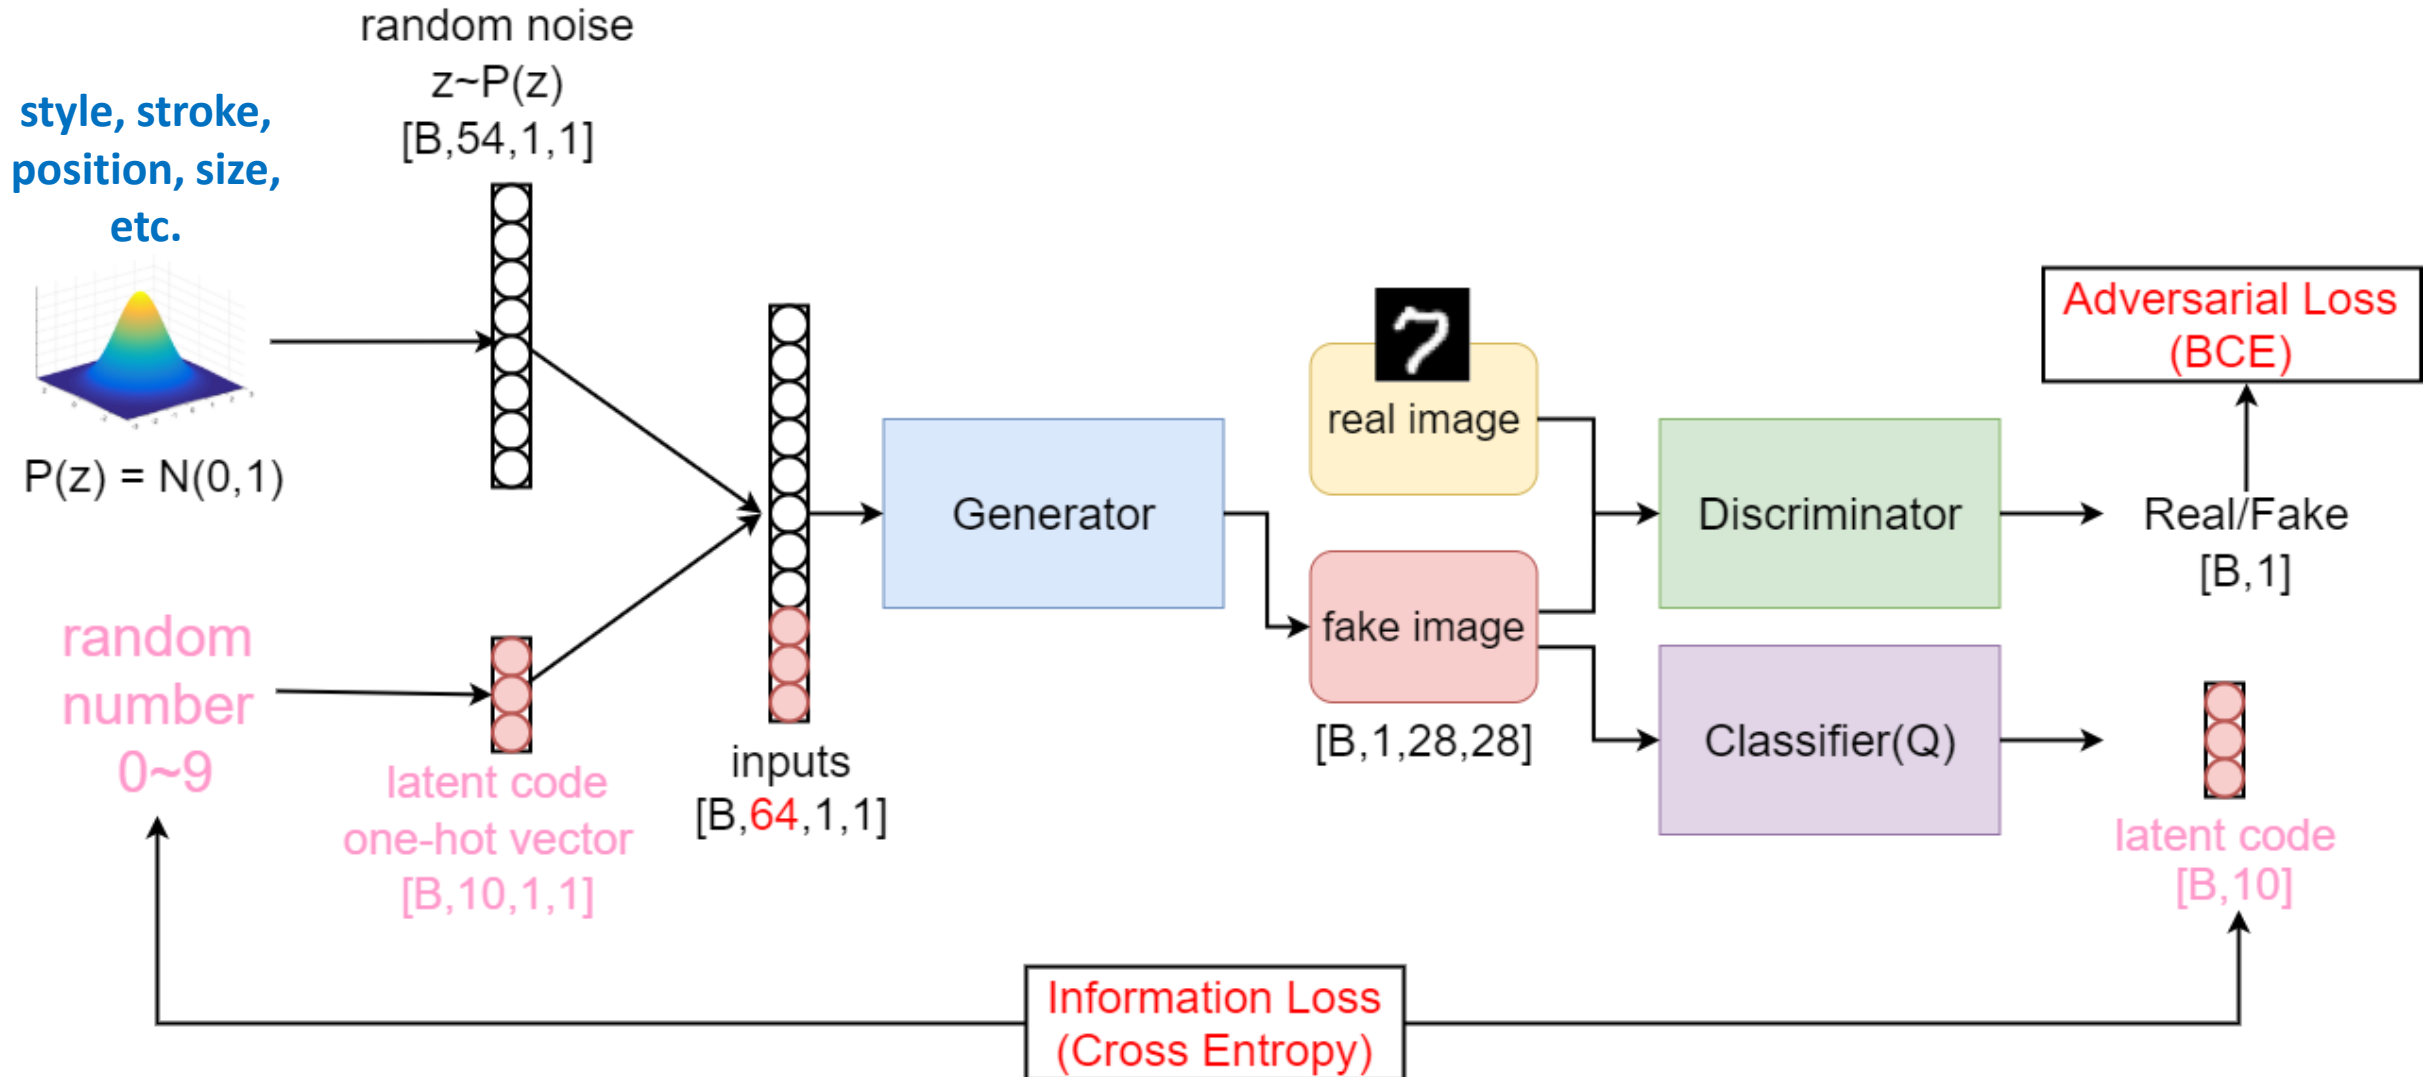
        # Discriminator branch
        self.D = nn.Sequential(
            #####
            # To Do
            #####
        )

        # Info branch
        self.Q = nn.Sequential(
            #####
            # To Do
            #####
        )

    def forward(self, x):
        h = self.main(x)
        real_or_fake = self.D(h).view(-1,1)
        info = self.Q(h).squeeze()

        return real_or_fake, info
```


InfoGAN-MNIST digits generation



Lab Requirement

- Modify the DCGAN architecture to InfoGAN.
 - Implement the model architecture in `model.py`
- Implement training procedure
 - Implement the training procedure in `train.py`
 - Adopt traditional generator and discriminator loss.
 - Maximize the mutual information between generated images and discrete one-hot vector.
- Show the generated images of digits.
- Plot the generator loss, discriminator loss and information loss during training.
- Implement the evaluation procedure.

Modify model.py

```
Generator(  
    (main): Sequential(  
      (0): ConvTranspose2d(64, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)  
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
      (2): LeakyReLU(negative_slope=0.01)  
      (3): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(1, 1), bias=False)  
      (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
      (5): LeakyReLU(negative_slope=0.01)  
      (6): ConvTranspose2d(256, 128, kernel_size=(5, 5), stride=(2, 2), padding=(2, 2), bias=False)  
      (7): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
      (8): LeakyReLU(negative_slope=0.01)  
      (9): ConvTranspose2d(128, 64, kernel_size=(2, 2), stride=(2, 2), bias=False)  
      (10): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
      (11): LeakyReLU(negative_slope=0.01)  
      (12): ConvTranspose2d(64, 1, kernel_size=(2, 2), stride=(2, 2), bias=False)  
      (13): Sigmoid()  
    )  
)  
  
Discriminator(  
    (main): Sequential(  
      (0): Conv2d(1, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
      (1): LeakyReLU(negative_slope=0.1, inplace=True)  
      (2): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
      (3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
      (4): LeakyReLU(negative_slope=0.1, inplace=True)  
      (5): Conv2d(128, 256, kernel_size=(7, 7), stride=(1, 1), bias=False)  
      (6): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
      (7): LeakyReLU(negative_slope=0.1, inplace=True)  
    )  
    (D): Sequential(  
      (0): Conv2d(256, 1, kernel_size=(1, 1), stride=(1, 1))  
      (1): Sigmoid()  
    )  
)
```

```
class Generator(nn.Module):  
    def __init__(self, in_dim=64):  
        super(Generator, self).__init__()  
  
        # Generator  
        self.main = nn.Sequential(  
            #####  
            # To Do  
            #####  
        )
```

```
class Discriminator(nn.Module):  
    def __init__(self):  
        super(Discriminator, self).__init__()  
  
        # shared layer of discriminator  
        self.main = nn.Sequential(  
            #####  
            # To Do  
            #####  
        )  
  
        # Discriminator branch  
        self.D = nn.Sequential(  
            #####  
            # To Do  
            #####  
        )  
  
        # Info branch  
        self.Q = nn.Sequential(  
            #####  
            # To Do  
            #####  
        )
```

Training procedure

- Update D network
 - Clear gradient
 - Real Loss
 - Input real data into D to get prediction
 - Calculate $\text{real_loss} = -\log(D(x))$ and backward
 - Fake Loss
 - Sample noise and generate one-hot vector, and concatenate them as the generator input
 - Pass the generator input into Generator to get fake data, and input fake data into D to get prediction
 - Calculate $\text{fake_loss} = -\log(1 - D(G(z)))$ and backward
- Update parameters

Training procedure

- Update G and Q network
 - Clear gradient
 - Generator Loss
 - Input fake data into D again to get prediction and latent vector
 - Calculate $G_loss = -\log(D(G(z)))$
 - Information Loss
 - Calculate $info_loss = \max(L_I(G, Q)) = \min(-c \cdot \log(Q(G(z, c))))$
 - Total Loss
 - Calculate $loss = G_loss + \lambda \cdot info_loss$ and backward
 - Update parameters


Hyper-parameters

- noise size = 54
- latent code size = 10
- total epochs = 50
- optimizer : Adam
- learning rate for G and Q: $1e^{-3}$
- learning rate for D: $2e^{-4}$
- Info loss weight: 0.25

You can adjust the hyper-parameters according to your own ideas.

Expected output

- 10 numbers (latent vectors) with 10 different styles (noises).



Demo

- Please **implement the evaluation procedure**
 - Load pretrained model weights

```
# load model
checkpoint = torch.load(os.path.join(opt.output_str, 'checkpoint/model_{}.pt'.format(epoch)))
netG = checkpoint['netG']
netD = checkpoint['netD']
```

- Generate output with 10 numbers x 10 styles

```
# generate data for evaluation procedure
fix_noise = torch.randn(10, opt.dim_noise, 1, 1, device=device).repeat_interleave(10, dim=0)
y = torch.arange(10).repeat(10)
y_cat = to_categorical(y, opt.dim_dis_latent)
val_data_inputs = torch.cat([fix_noise, y_cat], dim=1)
```

- While demo, you should run the evaluation and show your result to TA

Scoring Criteria

- Report (50%)
 - A. Introduction (10%)
 - B. Experiment Setups (20%)
 - How did you implement Info GAN
 - Adversarial loss
 - Maximizing mutual information
 - Which loss function of generator (refer to page.6) did you use? What's the difference?
 - C. Results and discussion (20%)
 - Results of your samples
 - Training loss curves

Scoring Criteria

- Demo (50%)
 - Generate 10 numbers with 10 different styles (noises). A column is correct if there are ≥ 7 correct numbers in that column, and your total score is decided according to the total number of correct columns. (30%)

• 10 correct columns	-----	100%
• 9 correct columns	-----	90%
• ≥ 7 correct columns	-----	80%
• ≥ 5 correct columns	-----	60%
• Otherwise	-----	0%
 - Questions (20%)

